

Genuine Semantic Publishing

[Incomplete Draft, 16 April 2016]

Tobias Kuhn and Michel Dumontier

Abstract—In the last years, many approaches and systems have been presented for what has been called *semantic publishing*. Closer inspection, however, reveals that these approaches are mostly not about publishing semantic representations, as the name seems to suggest. Rather, most approaches take the processes and outcomes of the current narrative-based publishing system for granted and only work with the already published papers. This includes semantic annotations, semantic interlinking, semantic integration, and semantic discovery, but with the semantics coming into play only after the publication of the original article. While these are interesting approaches, they fall short of providing a vision to transcend the current publishing paradigm. We argue for taking the term *semantic publishing* seriously and work towards a vision of *genuine* semantic publishing, where computational tools and algorithms can help us with dealing with the wealth of human knowledge by letting researchers capture their research results with formal semantics from the start. We argue that genuine semantic publications should come with formal semantics as an integral and primary component at the time of publication, that these representations should be complete in essence, in the sense that they cover the main results, that they should be authentic in the sense that they originate from the authors, and that they should be fine-grained and light-weight for optimized re-usability and minimized publication overhead. This paper is in fact not just advocating our concept, but is itself a *genuine semantic publication*, thereby demonstrating and illustrating our points.

I. INTRODUCTION

MANY scholars have pointed out that the classical way of publishing scientific articles is ill-suited to deal with the rapid growth of both, volume and complexity, of scientific contributions [1], [2]. To overcome these problems, next generation scientific publishing [3] has to respond to the increasing importance of datasets and software, and needs to provide methods to automatically organize and aggregate reported scientific findings. Perhaps the most important shortcoming of the current publication system is that scientific papers do not come with formal semantics that could be processed, aggregated, and interpreted in an automated fashion.

Semantic publishing [4]–[6] is a general approach to tackle this problem using the concepts and tools of the Semantic Web and related fields. This idea was basically born together with the idea of the Semantic Web itself. In 2001, Tim Berners-Lee and James Hendler sketched how they expect researchers in the future to produce machine-readable descriptions of their experiments and findings, in the form of mark-up of their research papers or as independent representations made public

on the web [7]. Unfortunately, subsequent work has deviated from this general proposal.

The topic of semantic publishing has received considerable attention during the last few years, with workshop series on Semantic Publishing (SePublica)¹ and Linked Science (LISC)² and two Semantic Publishing Challenges at ESWC conferences [8], [9]. However, as we argue below, these approaches interpret the term *semantic publishing* in a non-intuitive way: Instead of changing the publishing process, they mostly take existing classical publications as their starting point and simply apply semantic technologies on them, mostly without touching the publishing process or the object that is being published. This leads them to propose solutions that are quite timid and conventional, and fall short of providing a vision for the long-term future.

II. SEMANTIC PUBLISHING

Semantic publishing has been defined as “anything that enhances the meaning of a published journal article, facilitates its automated discovery, enables its linking to semantically related articles, provides access to data within the article in actionable form, or facilitates integration of data between papers” [5]. While this definition reflects what has been referred to as “semantic publishing” in the recent literature, we think that it is in one way too restrictive and in another way too inclusive, if we are serious about applying the ideas of the Semantic Web to scientific publishing.

The definition above is too restrictive in the sense that semantic publications according to this definition are required to accompany a “journal article” or a “paper.” An entity that contains only a semantic representation of a scientific result could not be considered a semantic publication unless it comes attached to a classical scientific article. On the other hand the definition is too inclusive, in our view, in the sense that it covers very shallow approaches that add little (if anything) to established approaches of publishing. For example, letting authors choose keywords from standardized vocabularies for their paper (as many journals do) in fact “enables its linking to semantically related articles,” and therefore by the definition above makes it a semantic publication. As another example, a semantic annotation performed by a third party on an article “enhances the meaning of a published journal article” and therefore would have to be called a semantic publication, even if the semantic annotation itself is not even made public. In the existing literature, the term “semantic publishing” seems in general to be interpreted as “adding semantics to something that is published” instead of the more intuitive readings of

Tobias Kuhn is with the Department of Computer Science at VU University Amsterdam, Netherlands; e-mail: t.kuhn@vu.nl.

Michel Dumontier is with the Stanford Center for Biomedical Informatics Research, Stanford University, CA, USA; e-mail: michel.dumontier@stanford.edu.

¹<http://sepublica.mywikipaper.org>

²<http://linkedscience.org/events/lisc2014/>




Fig. 1. The concept of *genuine semantic publishing* compared to what has been called *semantic publishing*, explained by an analogy where scientific papers are represented by boxes and formal semantics by flowers.

“publishing something that is semantic” or “publishing in a semantic manner.” As a side note, we would like to stress that we use the term *semantic* in its narrow technical sense that is common in the research field of the Semantic Web: By *semantic* we mean “with the meaning represented in a formal logic-based language” as compared to the broader original interpretation where *semantic* simply means “with respect to meaning.”

We argue here for a definition of *semantic publishing* that is broader in the sense that no narrative article needs to be present, that is at the same time narrower in the sense that the semantic representation has to be a first-class object created and published by the authors, and that is moreover consistent with the intuitive meaning of the term. We propose the following definition: *Genuine semantic publishing* occurs when a work is published that includes authentic and fine-grained representations of its content and meta data in a machine-interpretable semantic notation using shared conventions, and when these semantic representations are complete in essence and a primary component of the published entity. We explain in more detail below what we mean by *authentic*, *fine-grained*, *primary component*, and *complete in essence*.

By using a simple analogy, we can see classical papers as boxes that are closed and hard to access from the point of view of automated techniques. This situation is shown on the left hand side of Figure 1. Existing approaches to what has been called semantic publishing merely adorn this box with formal semantics — represented by flowers in the picture — but leaving it closed. This adornment is very useful, to be sure, but it does not reach to the main content of the box. By only looking at the formal semantics, one can possibly find out the topic of the paper but not what the paper is actually claiming. The main message is missed. Moreover, the adornment is often attached at a later point, after the box has been shipped so to say, and is therefore not a proper part of it. Speaking in terms

of this metaphorical image, we argue that we should open the box and let semantics bloom right from the inside. We should represent the paper’s main message with formal semantics. As we see on the right hand side of the figure, this metaphorically turns the box into a flowerpot. Now semantics is the main content, and the scientific paper has become a container for semantics instead of a closed box with a secondary usage as a pinboard for semantics annotations.

It seems to be a common unquestioned assumption that the semantic representation of knowledge has to start from a textual representation, and therefore writing a statement down in natural language always needs to be the first step. For example, we can read in a paper on semantic publishing that “learning how the brain creates and decodes meaning from text is essential if we are to provide better tools for scientific inquiry” and that we need to “train computers to help us read scientific text” [6]. While these are certainly interesting and important problems, it is not obvious why they are essential if we take the approach of semantic publishing seriously, i.e. if we ensure that the published artifacts come with semantic representations from the start. There is no law that research findings can only be formalized after they have been expressed in a narrative text. It can very well be the other way round, such as a researcher writing a narrative text verbalizing existing formal statements she has come up with. More likely, these two will go hand in hand in an iterative process, but this does not require computers to help us read scientific text. However, many articles in the area of semantic publishing make in fact this implicit assumption, as exemplified by papers presented at semantic publishing workshops claiming that “annotations on all levels pave the way for shared knowledge understanding” [10] and that “semantic publishing [...] can be defined as the activity of enhancing a document” [11], among many others (e.g. [12]–[14]). The entire approach of *semantic annotation* is based on this text-first paradigm, which is a perfectly valid

approach but is surely not the only possibility.

We get a similar picture if we look at the Semantic Publishing Challenge held at the Semantic Web conferences ESWC 2014 and 2015 [8], [9]. There were three tasks defined for each of these two challenges, but none of them actually deals with publishing. Instead they are about automatically extracting and interlinking semantic data from existing publications. Only the “in-use” task of the first challenge was general enough to not exclude publishing (“showcase the potential of Semantic Web technology for enhancing and assessing the quality of scientific production”), but it did not specifically mention the publishing process either. Unsurprisingly then, the approaches presented at these challenges deal with extraction from and annotation of articles that are already written and published (with the only exception being a paper introducing a publishing platform for Research Objects [15]).

III. RELATED WORK

Despite the prevalence of approaches that deal exclusively with semantic annotations and semantic interlinking of already published articles, there are a number of existing approaches where the artifacts to be published include from the start semantic representations that originate from the researchers themselves. They cover different aspects that we consider important for genuine semantic publishing. These approaches include Research Objects, executable papers, scholarly HTML, Structured Digital Abstracts, Micropublications, and Nanopublications.

Research Objects³ [1], [16] are a proposal to extend the notion of a scientific publication beyond a narrative article to include additional resources like data, meta data, code, presentation slides, log files, and workflow definitions. In order to improve reusability and reproducibility of scientific findings, these different resources are interlinked and packaged together in zip files called Research Objects. These objects are therefore at a higher level of composition than classical articles in the sense that they can contain articles and other resources, and we can say that their publishing unit is on the level of a *project* rather than the level of an article. The language RDF with its formal semantics is used as the underlying notation. The interpretation of a Research Object’s content, however, would typically be performed by a human user, at least according to the main motivating scenario [1]. The authors claim that Research Objects can also include formally represented hypotheses and high-level findings, but no specifics are given on how this can be achieved. Later work by other researchers applied the proposed concepts and used embedded nanopublications (see below) to represent such higher-level claims and hypotheses [17].

From a different angle, a number of approaches have been proposed for what has been called *executable papers* (e.g. [18]–[20]), in particular in response to the Executable Paper Grand Challenge⁴ organized by Elsevier in 2011. These approaches follow the same basic idea as the IPython Notebook⁵, which has recently become very popular, to interweave

narrative text with program code that can be executed in an interactive fashion by the reader to dynamically generate results in the form of tables and diagrams, and thereby to reproduce, verify, and explore the results. Many of the proposed solutions do not explicitly mention semantic interlinking or standardized vocabularies, and they are therefore a borderline case of semantic publishing. However, one can argue that source code is a kind of semantic representation, even though with procedural rather than declarative semantics. Naturally, these approaches focus on algorithms, their source code, and the corresponding input and output data, and do not target the formal representation of meta data or high-level claims and hypotheses.

Yet another angle is taken by approaches on scholarly HTML⁶ that propose HTML-based replacements for today’s role of PDF as universal file format for scientific articles. Many of these approaches do not explicitly focus on formal semantics, but — in contrast to PDF — semantic markup comes very naturally with HTML via technologies such as RDFa. Initiatives like Linked Research [21] propose general principles and directions of using HTML for scholarly communication to enable, among other things, the publication of semantic representations. Specific HTML-based formats such as RASH [22] provide specific guidelines and tools of how that can be achieved. The semantic markup of such articles is mostly on the meta level, for example about the authors, the structure of the article, and the references to other articles. Furthermore, as the term *markup* indicates, these semantic representations are mostly not independent statements but tightly linked to the narrative text. The possible use of RDFa to formally represent not just meta data but also high-level claims, hypotheses, and arguments is sometimes proposed, but no concrete solutions are presented.

Structured Digital Abstracts [23], [24], having been proposed almost ten years ago, are probably the oldest approach of this kind. They were revolutionary in the sense that they proposed to let authors capture the main claims of their scientific contributions, such as newly discovered protein-protein interactions, with machine-interpretable semantics. Even though these abstracts are attached to narrative articles, the formally represented findings can be processed and interpreted independently from the narrative text. We proposed a similar approach in previous work with abstracts in controlled natural language [25].

Micropublications [26] are a further approach, which puts the emphasis on the structure and interrelation of scientific arguments and their underlying pieces of evidence. The authors stress that the network of arguments is an essential part of science, of which claims and hypotheses are necessary but not sufficient ingredients. They argue that formal representations of scientific claims are often not practically feasible, whereas the structure among them can be captured more easily and is moreover more important and more valuable to help scientists with computer-aided knowledge management.

Nanopublications are an approach to use the RDF language to represent “the smallest unit of publication” [27] or

³<http://www.researchobject.org>

⁴<http://www.executablepapers.com>

⁵<http://ipython.org/notebook.html>

⁶<https://www.w3.org/community/scholarlyhtml/>

		Research Objects	executable papers	scholarly HTML	Structured Digital Abstracts	Micro-publications	Nano-publications
types of formal semantic representations:	– meta data	✓		✓	✓	✓	✓
	– domain (input/output) data	✓	✓				✓
	– program code	✓	✓				
	– arguments / discourse			(✓)		✓	(✓)
publishing unit:	– high-level claims	(✓)		(✓)	✓		✓
	– project level	✓					
	– article level		✓	✓	✓	✓	
	– statement level						✓
necessary components:	– formal semantic data	✓	✓	(✓)	✓	✓	✓
	– narrative text		✓	✓	✓	✓	

TABLE I
PROPERTIES OF EXISTING APPROACHES ON THE PUBLICATION OF SCIENTIFIC ARTIFACTS CONTAINING SEMANTIC REPRESENTATIONS

“core scientific statements with associated context” [28]. This statement-level approach is therefore at a more granular level than most other approaches whose unit of publication is at the article level. Nanopublications consist of three parts, each represented in RDF: an assertion containing the actual content in the form of an atomic small piece of knowledge (e.g. a scientific claim, or a data entry), a provenance part containing meta data about the origin and context of the assertion (e.g. how it was measured), and a publication information part with meta data about the nanopublication as a whole (e.g. when and by whom it was created). Even though the details of their integration into the scientific publishing workflow have remained largely unspecified, nanopublications have received considerable attention during the last few years, with several large dataset having been published in this format [29]–[31]. In our own previous work, we have proposed a preliminary general approach of representing within nanopublications the structure among informal claims and hypotheses, which are themselves not necessarily formally represented [32], thereby addressing some of the points raised by micropublications.

Table I shows a comparison and summary of the different types of approaches introduced above. We see that they cover different types of semantic representations. We also observe from the table that none of the existing approaches covers all aspects, but a complete coverage could be achieved by a combination of them (and there don’t seem to be any fundamental incompatibilities). We will argue below that high-level claims and arguments are the most important parts, because they are — and have always been — the core contributions of a scientific work.

Executable papers, Structured Digital Abstracts, and Micropublications stick to the article as their unit of publication, whereas Research Objects operate at a higher level (at what we call the project level) and nanopublications at a lower one (statement level). We argue here that it is important to cover all these levels, which means that it has to be possible to publish something as small as an individual statement (i.e. an individual claim, hypothesis, data entry, or link between entities). Once this lowest level is covered, it is then relatively straightforward to compose larger objects out of the low-level ones, like the approach mentioned above of using nanopublications within Research Objects [17].

All these approaches, except some flavors of scholarly HTML, mandate that formal semantic data are part of the

published entity, but three out of the five approaches also require or assume that a narrative text accompanies the data. Not coincidentally, these are also the approaches that work on the article level, and therefore stick to the classical unit of publication. We will argue below that narrative text necessarily remains an important part of scientific discourse and communication, but it also has to be possible to publish data that is self-explanatory due to its formal semantics without the need for a narrative.

To summarize, there are a number of existing approaches that cover important aspects of what we think deserves the term *semantic publishing* under its intuitive interpretation, but they are normally overshadowed by the more conventional approaches of semantic annotation and semantic interlinking introduced above. These conventional approaches, in fact, advocate the term *semantic publishing* more visibly, whereas approaches like Research Objects often don’t use the term at all. This situation can lead to a general impression that the conventional approaches are the best we can do to apply the ideas of the Semantic Web to scientific publishing. It is our goal to convince the readers of this paper otherwise.

To conclude our discussion of related work, we would like to point out that a large number of general technologies have been developed in the last years that can serve as the basis for approaches on genuine semantic publishing. They come in the form of data formats, ontologies, and software tools. The Semantic Publishing and Referencing Ontologies (SPAR)⁷ [33], [34], for example, are a highly valuable suite of ontologies on the domain of scientific publishing. Other examples include the Annotation Ontology [35] that defines how to connect (scientific) text to the respective formal representations, the PROV Ontology [36] to model provenance of digital object, and the Linked Science Core Vocabulary⁸ that covers relations between entities in the scientific workflows such as researchers, data, publications, and methods. There is also much existing work on modeling arguments and discourse with Semantic Web technology [37], including approaches on formalizing uncertainty in scientific arguments [38].

IV. GENUINE SEMANTIC PUBLISHING

As we have shown above, most approaches that go under the label *semantic publishing* are not actually about publishing,

⁷<http://sempublishing.sourceforge.net>

⁸<http://linkedscience.org/lsc/ns/>

and the approaches that *do* target the publication of semantic representations cover different aspects thereof that only partly overlap. We therefore think that there is a need for clear criteria of *genuine* semantic publishing. As we started to argue above, important aspects of semantic publishing concern the authoritativeness and completeness in essence of semantic representations, as well as their status in relation to narrative articles and their granularity level.

The first aspect we would like to discuss here is what we call *completeness in essence* of semantic representations with respect to the entity to be published. A representation is complete in essence with respect to a work if it covers (at least) the essence of the work. The essence of a work is its main message, which for scientific articles normally consists of the main claims, findings, and arguments. A semantic representation may not cover all aspects discussed and described throughout a scientific work, but for it to be complete in essence it has to cover the main points: If you had to summarize a paper in one sentence, the content of this sentence has to be present in the semantic representation too. In other words, it is not sufficient to focus on what is *easy* to represent; we have to focus on what is *important* to represent. One can also see it as a kind of democratization process of making automated agents first-class citizens: English speaking agents (e.g. human researchers) get the main content of the work in their English based representation (i.e. the narrative text); so RDF speaking agents (e.g. Linked Data aware software) should also get the main content of the work in their RDF based representation. As we will see below, this perspective aligns very well with the well-established Web technique of content negotiation.

Another important aspect is the authoritativeness of the source of the semantic representations, which determines the authenticity of these representations. Semantic representations can only be considered authentic if they originate from an agent that is authoritative in the given situation. In the case of the publication of a scientific result, the only authoritative source are the researchers (which are called *authors* in this context). In other words, semantic representations of scientific results are only authentic if they are provided by the researchers themselves. It has long been known in the area of knowledge engineering that the process of formalizing expert knowledge is not merely a process of “transferring” or “converting” knowledge from existing representations inside the heads of experts to formal representations of a form that can be stored in a knowledge base. Rather it has to be seen as a creative *modeling process* [39], [40] where formal structures are generated that existed only in an incomplete, implicit, and unconscious form in the experts’ heads. Explaining a result in a narrative is simpler than formally modeling it, in the sense that natural language allows the writer to remain vague and even ambiguous. Accurately modeling knowledge only from a narrative text without a further connection to the authoritative source is therefore in general not just difficult but strictly impossible. Genuine semantic publishing requires the authors of scientific results to perform the modeling task themselves, because they are — by definition — the only authoritative source. We claim that — contrary to many existing approaches — we should *not* try to relieve the authors

of this burden (though we should of course try to help them). Otherwise, the semantic representations cannot be considered authoritative and should therefore not be considered part of the publication’s content (unless the person who produces the formal representations becomes a co-author), and we would end up with a situation where semantic representations are disconnected from what is published, which is against the essence of the semantic publishing idea. As Tim Berners-Lee and James Hendler made clear when the Semantic Web was just about to come into existence as a research field, it “involves asking people to make some extra effort” [7]. In the case of scientific publishing, it involves asking authors to make the extra effort of providing formal semantic representations of their findings.

[...]

Based on the arguments presented above, we claim that *genuine* semantic publishing needs to comply with the following criteria:

- 1) A scientific work needs to come with formal representations that are semantic, in the sense that they are not just machine processable but *machine interpretable*, and that are linked so they add to the existing formal body of knowledge and can be used to answer questions or aggregate results.
- 2) These semantic representations need to be *complete in essence* (though possibly underspecified) in the sense that both domain data and meta data are included, as well as links between the two, and that they cover (at least) the core of the main results of the given work.
- 3) They need to be *authentic* in the sense that the respective authoritative persons create or approve the semantic representations. Domain data can only come directly from the researchers, and meta data has to come from the people responsible for the form of the published work (i.e. the researchers and/or the editors).
- 4) The semantic representations need to be a *primary* component of the published work, made available together with everything else at the time of publication. They must furthermore have an independent existence in their own right and not merely be appended or attached to the main entity as an appendix or as annotations.
- 5) The semantic representations and their containers need to be *fine-grained and light-weight*. Even though such semantic representations might often be published in large chunks, the publication of minimal additions and corrections need to be possible without a large overhead.

Most, maybe all existing approaches on what has been called semantic publishing comply with the first criterion, but only a few of them propose or support representations that comply with the other criteria of being complete in essence, authentic, primary, and fine-grained. Below we illustrate how existing technologies can be combined and applied to achieve such genuine semantic publishing.

V. GENUINE SEMANTIC PUBLISHING IN ACTION

It turns out that all the technologies needed for applying genuine semantic publishing are already available and most

of them are very mature and reliable. There are no technical obstacles preventing us from releasing our results from today on as genuine semantic publications (though of course more work is needed on nice and intuitive end-user interfaces to make this process as easy as possible).

The paper that you are reading is in fact a genuine semantic publication. It has different representations for different types of usage. [...] We use CiTO [33] and SKOS [41]. [...]

We have to formally represent the main content of the paper. Our paper's main message is the advocacy of the new concept of genuine semantic publishing:

```
p:paper
  cito:describes p:GenuineSemanticPublishing ;
  cito:supports p:GenuineSemanticPublishing .
```

There is to our knowledge no existing ontology that would exactly capture the relation of a publication *advocating* a given concept, but the combination of the two relations *describes* and *supports* from the CiTO ontology comes close. We as authors should of course say a bit more about this new concept, most importantly that it is related to the existing concept of semantic publishing:

```
p:GenuineSemanticPublishing skos:related
  dbpedia:Semantic_publishing .
```

And we can express our critical position on that concept:

```
p:paper cito:critiques
  dbpedia:Semantic_publishing .
```

Next we can formally represent the five criteria based on which we define our new concept:

```
p:GenuineSemanticPublishing skos:definition
  p:GenuineSemanticPublishingCriteria .

p:GenuineSemanticPublishingCriteria dct:hasPart
  p:GenuineSemanticPublishingCriterion1 .

p:GenuineSemanticPublishingCriterion1
  dct:title "First criterion genuine semantic
    publishing: machine interpretability" ,
  dct:description "A scientific work needs to
    come with formal representations that are
    semantic, in the sense that ..." .

p:GenuineSemanticPublishingCriteria dct:hasPart
  p:GenuineSemanticPublishingCriterion2 .
```

...

We can try to capture part of the content of these criteria in RDF as well, but at some point we have to stop and be content with an informal description in natural language (at the latest when we hit the symbol grounding problem). However, we believe that it is always possible to find a formal representation of the main content at the highest level (such as introducing and advocating a new concept), even though we will mostly not be able to provide a complete formal definition. In this sense such a representation is *complete in essence, but underspecified*.

At this point, the benefits of such a representation of the main message of a paper might not seem obvious. One of the main benefits comes when *subsequent* papers or nanopublications start referring to these formal representations. For example, a subsequent paper might propose the concept of “fantastic publishing” that includes our criteria 1 to 4, but criticizes number 5 and suggests to replace it with a different one:

```
p2:anotherPaper
  cito:describes p2:FantasticPublishing ;
  cito:supports p2:FantasticPublishing .

p2:FantasticPublishing skos:related
  p:GenuineSemanticPublishing .

p2:anotherPaper cito:critiques
  p:GenuineSemanticPublishingCriterion5 .

p2:FantasticPublishing skos:definition
  p:FantasticPublishingCriteria .

p2:FantasticPublishingCriteria dct:hasPart
  p:GenuineSemanticPublishingCriterion1 ,
  p:GenuineSemanticPublishingCriterion2 ,
  p:GenuineSemanticPublishingCriterion3 ,
  p:GenuineSemanticPublishingCriterion4 ,
  p2:FantasticPublishingCriterion .

p2:FantasticPublishingCriterion
  dct:title "Fantastic criterion for fantastic
    publishing" ,
  dct:description "...".
```

This admittedly silly example shows how we can formally capture the high-level relation of papers' content. [...]

We can use HTML as a format to simply link to different (classical and semantic) representations and at the same time to provide a simple landing page. With just a few lines of HTML code, we can define a canonical URL and some minimal metadata (such as title and authors of the work; more metadata is available in the actual representations):

```
<!DOCTYPE html>
<html>
<head>
<link rel="canonical"
  href="http://www.tkuhn.org/pub/sempub/">
<title>Genuine Semantic Publishing</title>
</head>
<body>
<h1>Genuine Semantic Publishing</h1>
<p>
by
<a href="http://orcid.org/0000-0002-1267-0234"
  rel="author">Tobias Kuhn</a>
and
<a href="http://orcid.org/0000-0003-4727-9435"
  rel="author">Michel Dumontier</a>
</p>
```

And then we can link to different representations of the content of the given work:

```
<p>Content:</p>
<ul>
<li><a rel="alternate" href="sempub.pdf"
  type="application/pdf">as PDF</a></li>
<li><a rel="alternate" href="sempub.dokieli.html"
  type="text/html">as HTML/Dokieli</a></li>
<li><a rel="alternate" href="sempub.rash.html"
  type="text/html">as HTML/RASH</a></li>
<li><a rel="alternate" href="sempub.ttl"
  type="text/turtle">as RDF/Turtle</a></li>
<li><a rel="alternate" href="sempub.trig"
  type="application/trig">as RDF/Trig</a></li>
</ul>
</body>
</html>
```

Figure 2 shows what such a minimal landing page looks like in a browser.

It is important to stress that these are different representations of the *same* work that are each *complete in essence*. In other words, the main content of the work (such as the main



Fig. 2. The landing page pointing to different versions of the work.

conclusions in the case of a scientific paper) is represented in *all* available representations. In this sense, these different representations are very different from the supplementary material that is often provided for scientific articles.

[...]

VI. DISCUSSION AND CONCLUSIONS

[...]

ACKNOWLEDGMENT

We would like to thank Silvio Peroni for his advice on ontologies with respect to scholarly communication. [...] Figure 1 was designed by Germán Barboza, from Cordero Producciones.⁹

REFERENCES

- [1] S. Bechhofer, I. Buchan, D. De Roure, P. Missier, J. Ainsworth, J. Bhagat, P. Couch, D. Cruickshank, M. Delderfield, I. Dunlop *et al.*, "Why linked data is not enough for scientists," *Future Generation Computer Systems*, vol. 29, no. 2, pp. 599–611, 2013.
- [2] J. Priem, "Scholarship: Beyond the paper," *Nature*, vol. 495, no. 7442, pp. 437–440, 2013.
- [3] T. Clark, "Next generation scientific publishing and the web of data," *Semantic Web*, vol. 5, no. 4, pp. 257–259, 2014.
- [4] D. Shotton, K. Portwin, G. Klyne, and A. Miles, "Adventures in semantic publishing: exemplar semantic enhancements of a research article," *PLoS computational biology*, vol. 5, no. 4, p. e1000361, 2009.
- [5] D. Shotton, "Semantic publishing: the coming revolution in scientific journal publishing," *Learned Publishing*, vol. 22, no. 2, pp. 85–94, 2009.
- [6] A. de Waard, "From proteins to fairytales: directions in semantic publishing," *Intelligent Systems, IEEE*, vol. 25, no. 2, pp. 83–88, 2010.
- [7] T. Berners-Lee and J. Hendler, "Publishing on the semantic web," *Nature*, vol. 410, no. 6832, pp. 1023–1024, 2001.
- [8] C. Lange and A. Di Iorio, "Semantic publishing challenge—assessing the quality of scientific output," in *Semantic Web Evaluation Challenge*. Springer, 2014, pp. 61–76.
- [9] A. Di Iorio, C. Lange, A. Dimou, and S. Vahdati, "Semantic publishing challenge—assessing the quality of scientific output by information extraction and interlinking," in *Semantic Web Evaluation Challenges*. Springer, 2015, pp. 65–80.
- [10] P. Smrz and J. Dytrych, "Towards new scholarly communication: A case study of the 4a framework," in *First Workshop on Semantic Publication (SePublica 2011)*. CEUR-WS, 2011.
- [11] A. Ruiz-Iniesta and O. Corcho, "A review of ontologies for describing scholarly and scientific documents," in *4th Workshop on Semantic Publishing (SePublica 2014)*. CEUR-WS, 2014.
- [12] A. C. S. Croset, S. Kafkas, M. Liakata, and A. Oellrich, "Exploring the generation and integration of publishable scientific facts using the concept of nano-publications," in *First Workshop on Semantic Publication (SePublica 2011)*. CEUR-WS, 2011.
- [13] C. H. Marcondes, "A semantic model for scholarly electronic publishing," in *First Workshop on Semantic Publication (SePublica 2011)*. CEUR-WS, 2011.
- [14] A. Di Iorio, S. Peroni, F. Vitali, and J. Zingoni, "Semantic lenses to bring digital and semantic publishing together," 2014.
- [15] R. Palma, P. Hołubowicz, O. Corcho, J. M. Gómez-Pérez, and C. Mazurek, "ROHub — a digital library of research objects supporting scientists towards reproducible science," in *Semantic Web Evaluation Challenge*. Springer, 2014, pp. 77–82.
- [16] S. Bechhofer, D. De Roure, M. Gamble, C. Goble, and I. Buchan, "Research objects: Towards exchange and reuse of digital knowledge," *The Future of the Web for Collaborative Science*, 2010.
- [17] A. González-Beltrán, P. Li, J. Zhao, M. S. Avila-Garcia, M. Roos, M. Thompson, E. van der Horst, R. Kaliyaperumal, R. Luo, T.-L. Lee *et al.*, "From peer-reviewed to peer-reproduced in scholarly publishing: the complementary roles of data models and workflows in bioinformatics," *PLOS one*, vol. 10, no. 7, p. e0127612, 2015.
- [18] P. Nowakowski, E. Ciepiela, D. Hareźlak, J. Kocot, M. Kasztelnik, T. Bartyski, J. Meizner, G. Dyk, and M. Malawski, "The collage authoring environment," *Procedia Computer Science*, vol. 4, pp. 608–617, 2011.
- [19] P. Van Gorp and S. Mazanek, "Share: a web portal for creating and sharing executable research papers," *Procedia Computer Science*, vol. 4, pp. 589–597, 2011.
- [20] M. Kohlhasse, J. Corneli, C. David, D. Ginev, C. Jucovschi, A. Kohlhasse, C. Lange, B. Matican, S. Mirea, and V. Zholudev, "The planetary system: Web 3.0 & active documents for STEM," *Procedia Computer Science*, vol. 4, pp. 598–607, 2011.
- [21] S. Capadisli, R. Riedl, and S. Auer, "Enabling accessible knowledge," *CeDEM15: Conference for E-Democracy and Open Government*, p. 257, 2015. [Online]. Available: <http://csarven.ca/enabling-accessible-knowledge>
- [22] A. Di Iorio, A. G. Nuzzolese, F. Osborne, S. Peroni, F. Poggi, M. Smith, F. Vitali, and J. Zhao, "The RASH framework: enabling HTML + RDF submissions in scholarly venues," in *In Proceedings of the ISWC 2015 Posters & Demonstrations Track*, 2015. [Online]. Available: http://ceur-ws.org/Vol-1486/paper_72.pdf
- [23] M. R. Seringhaus and M. B. Gerstein, "Publishing perishing? towards tomorrow's information architecture," *BMC bioinformatics*, vol. 8, no. 1, p. 17, 2007.
- [24] A. Ceol, A. Chatr-Aryamontri, L. Licata, and G. Cesareni, "Linking entries in protein interaction database to structured text: the febs letters experiment," *FEBS letters*, vol. 582, no. 8, pp. 1171–1177, 2008.
- [25] T. Kuhn, L. Royer, N. E. Fuchs, and M. Schroeder, "Improving Text Mining with Controlled Natural Language: A Case Study for Protein Interactions," in *3rd International Workshop on Data Integration in the Life Sciences 2006 (DILS'06)*. Springer, 2006.
- [26] T. Clark, P. Ciccarese, and C. Goble, "Micropublications: a semantic model for claims, evidence, arguments and annotations in biomedical communications," *Journal of Biomedical Semantics*, vol. 5, no. 1, p. 28, 2014.
- [27] B. Mons, H. van Haagen, C. Chichester, J. T. den Dunnen, G. van Ommen, E. van Mulligen, B. Singh, R. Hooft, M. Roos, J. Hammond *et al.*, "The value of data," *Nature genetics*, vol. 43, no. 4, pp. 281–283, 2011.
- [28] P. Groth, A. Gibson, and J. Velterop, "The anatomy of a nano-publication," *Information Services and Use*, vol. 30, no. 1, pp. 51–56, 2010.
- [29] C. Chichester, P. Gaudet, O. Karch, P. Groth, L. Lane, A. Bairoch, B. Mons, and A. Loizou, "Querying nextprot nanopublications and their value for insights on sequence variants and tissue expression," *Web Semantics: Science, Services and Agents on the World Wide Web*, 2014.
- [30] J. M. Banda, T. Kuhn, N. H. Shah, and M. Dumontier, "Provenance-centered dataset of drug-drug interactions," in *Proceedings of the 14th International Semantic Web Conference (ISWC 2015)*. Springer, 2015, pp. 293–300.
- [31] N. Queralt-Rosinach, T. Kuhn, C. Chichester, M. Dumontier, F. Sanz, L. I. Furlong *et al.*, "Publishing disgenet as nanopublications," *Semantic Web-Interoperability, Usability, Applicability*, 2016.
- [32] T. Kuhn, P. E. Barbano, M. L. Nagy, and M. Krauthammer, "Broadening the scope of nanopublications," in *Proceedings of the 10th Extended Semantic Web Conference (ESWC 2013)*. Springer, 2013.

⁹<http://designers.designcrowd.com/designer/32548/cordero-producciones>

- [33] S. Peroni and D. Shotton, "FaBiO and CiTO: ontologies for describing bibliographic resources and citations," *Web Semantics: Science, Services and Agents on the World Wide Web*, vol. 17, pp. 33–43, 2012.
- [34] S. Peroni, D. Shotton, and F. Vitali, "Scholarly publishing and linked data: describing roles, statuses, temporal and contextual extents," in *Proceedings of the 8th International Conference on Semantic Systems*. ACM, 2012, pp. 9–16.
- [35] P. Ciccarese, M. Ocana, L. J. G. Castro, S. Das, and T. Clark, "An open annotation ontology for science on web 3.0," *Journal of biomedical semantics*, vol. 2, no. 2, p. 1, 2011.
- [36] T. Lebo, S. Sahoo, D. McGuinness, K. Belhajjame, J. Cheney, D. Corsar, D. Garijo, S. Soiland-Reyes, S. Zednik, and J. Zhao, "PROV-O: The PROV ontology," *W3C Recommendation*, 2013.
- [37] J. Schneider, T. Groza, and A. Passant, "A review of argumentation for the social semantic web," *Semantic Web*, vol. 4, no. 2, pp. 159–218, 2013.
- [38] A. De Waard and J. Schneider, "Formalising uncertainty: An ontology of reasoning, certainty and attribution (ORCA)," *Semantic Technologies Applied to Biomedical Informatics and Individualized Medicine*, 2012.
- [39] W. J. Clancey, "The knowledge level reinterpreted: Modeling how systems interact," *Machine Learning*, vol. 4, no. 3-4, pp. 285–291, 1989.
- [40] R. Studer, V. R. Benjamins, and D. Fensel, "Knowledge engineering: Principles and methods," *Data and Knowledge Engineering*, vol. 25, no. 1–2, pp. 161–197, March 1998.
- [41] A. Miles and S. Bechhofer, "SKOS simple knowledge organization system reference," *W3C recommendation*, vol. 18, p. W3C, 2009.